

QATIP Intermediate

AWS Lab 06

Advanced Terraform Expressions, Functions, and Regular Expressions

Contents

Lab Objectives	1
Challenge 1: Creating an S3 Bucket with Dynamic Naming and Tagging	2
Scenario	2
Goal: Define Terraform Configuration for S3 Bucket	2
Validation Checks	3
Run tests with invalid parameters:	3
Run tests with valid parameters	3
Hints.....	3
Solution	3
Challenge 2: Standardizing Resource Names Using Regular Expressions	3
Scenario	3
Goal: Load JSON File and Process Data using Regex.....	4
Hints.....	4
Solution	4

Lab Objectives

This multi-challenge lab integrates expressions, dynamic blocks, Terraform functions, and regular expressions to create flexible and reusable infrastructure configurations in AWS. By the end of this lab, you will:

- Implement conditional expressions to control resource creation.

- Leverage string, numeric, and date/time functions to standardize S3 bucket naming, tagging, and configuration.
- Utilize regular expressions (regex) to standardize inconsistent AWS resource names.
- Validate outputs using terraform console.

Before you begin

- Ensure you have completed Lab0 before attempting this lab.
- In the IDE terminal pane, enter the following command...
`cd ~/environment/aws-tf-int/labs/06`
- There is a sub-directory for each of the following challenges. Navigate to the appropriate folder as you attempt each challenge. Ensure all file changes and commands are executed in the appropriate challenge subdirectory

Challenge 1: Creating an S3 Bucket with Dynamic Naming and Tagging

Scenario

You are tasked with creating a reusable Terraform configuration that:

- Creates an S3 bucket with a name containing “bucket1” through to “bucket6”.
- Adds a mandatory **Environment** tag that is always uppercase and restricted to **DEV**, **PROD**, or **TEST**.
- Adds a second tag, **UpdatedOn**, with the creation date in the format **YYYY-MM-DD**.
- Uses Terraform variables with default values for the S3 bucket name and environment.
- Implements validation checks to ensure:
 - The S3 bucket name is within the specified range.
 - The environment value is restricted to DEV, PROD, or TEST.

Goal: Define Terraform Configuration for S3 Bucket

Updating the files provided in `/aws-tf-int/labs/06/challenge1`, attempt to create a deployment that complies with the given scenario.

Validation Checks

To validate your configuration, run **terraform plan**

Review the output to ensure that the S3 bucket name, Environment tag, and UpdatedOn tag meet the required constraints.

Run tests with invalid parameters:

```
terraform plan -var="s3_bucket_name=bucket8" -var="environment=Plan"
```

Expected error output:

Error: Bucket name must contain “bucket1” through “bucket6”.

Error: Environment must be one of DEV, PROD, or TEST.

Run tests with valid parameters

```
terraform plan -var="s3_bucket_name=bucket1" -var="environment=dev"
```

Expected success message:

The environment should be converted to uppercase (DEV).

Hints

- Use `format()` and `upper()` for consistency.
- Use `timestamp()` to capture creation date/time.
- Use `formatdate()` to generate the UpdatedOn tag.
- Use `contains()` to validate the S3 bucket name.

Solution

A proposed solution to this challenge can be found at [solutions/06/challenge1](#)

Challenge 2: Standardizing Resource Names Using Regular Expressions

Scenario

You’ve been assigned to the Cloud Infrastructure team at a company migrating virtual machines (VMs) from multiple departments to AWS. Each department has used its own naming convention, leading to inconsistent VM naming such as:

<code>`dev-finance.db01`</code>	<code>(stage-project.name)</code>
<code>`infra_prod_appserver`</code>	<code>(project_stage_name)</code>
<code>`sales-test_db02`</code>	<code>(project-stage_name)</code>
<code>`infra_prod_web01`</code>	<code>(project_stage_name)</code>
<code>`test.backup.storage03`</code>	<code>(stage.project.name)</code>

The business has mandated a standardized naming convention for all AWS EC2 instances: **[environment]-[project]-[name]** (e.g., *prod-finance-db01*)

Goal: Load JSON File and Process Data using Regex

Updating the files provided in `/aws-tf-int/labs/06/challenge2`, attempt to create a deployment that complies with the given scenario

- Use Terraform functions to process a JSON file containing inconsistent virtual machine names.
- Reconstruct the names in accordance with the mandated naming convention using Regex.
- Load the JSON file **resource-names.json** using `jsondecode()`
- Construct regex patterns dynamically based on known stages and projects
- Extract the environment, project, and machine name using `regex()`
- Generate new names in accordance with the naming convention
- Generate screen output listing the original names and the converted names:

```
+ "dev-finance.db01"      = "dev-finance-db01"
+ infra_prod_appserver    = "prod-infra-appserver"
+ infra_prod_web01        = "prod-infra-web01"
+ sales-test_db02         = "test-sales-db02"
+ "test.backup.storage03" = "test-backup-storage03"
```

- There is no requirement to create these machines at this stage

Hints

- Use `jsondecode(file("./resource-names.json"))` to read the JSON file
- Use `join("|", {stages/projects})` logic to create a regex pattern for stages and projects to find matches against
- Use `try(regex(pattern, name), "default_value")` to avoid errors on missing matches

Solution

A proposed solution for this challenge can be found at `solutions/06/challenge2`